

# Predicting Metro Interstate Traffic Volume: A Time Series Forecasting Approach with ARIMA

KJ MoChroi

Department of Data Science, Bellevue University

DSC680: Applied Data Science

Dr. Brett Werner

Spring 2023

Dataset: <https://archive.ics.uci.edu/ml/datasets/Metro+Interstate+Traffic+Volume>

## Import and View Data

```
In [1]: # Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [2]: # import dataset and preview
df = pd.read_csv("Potential_Datasets/Metro_Interstate_Traffic_Volume.csv.gz", compression='gzip')
df.head()
```


```
Out[2]:
```

	holiday	temp	rain_1h	snow_1h	clouds_all	weather_main	weather_description	date_time	traffic_v
0	None	288.28	0.0	0.0	40	Clouds	scattered clouds	2012-10-02 09:00:00	
1	None	289.36	0.0	0.0	75	Clouds	broken clouds	2012-10-02 10:00:00	
2	None	289.58	0.0	0.0	90	Clouds	overcast clouds	2012-10-02 11:00:00	
3	None	290.13	0.0	0.0	90	Clouds	overcast clouds	2012-10-02 12:00:00	
4	None	291.14	0.0	0.0	75	Clouds	broken clouds	2012-10-02 13:00:00	

```
In [3]: # Let's make sure 'date' is actually a date in pandas
df["date_time"] = pd.to_datetime(df["date_time"])
```


In [4]:

```
# plot the dataset
fig, ax = plt.subplots(figsize=(20, 6))
ax.plot(df["date_time"], df["traffic_volume"]);
```




In [5]:

```
# plot the dataset
fig, ax = plt.subplots(figsize=(20, 6))
ax.plot(df["date_time"][20000:40000], df["traffic_volume"][20000:40000]);
```




In [6]:

```
# plot the dataset
fig, ax = plt.subplots(figsize=(20, 6))
ax.plot(df["date_time"][20000:22000], df["traffic_volume"][20000:22000]);
```



In [7]:

```
# plot the dataset
fig, ax = plt.subplots(figsize=(20, 6))
ax.plot(df["date_time"][20000:20200], df["traffic_volume"][20000:20200]);
```



## Clean Dataset

```
In [8]: # drop data prior to 2015 - 07
df_complete = df[df["date_time"] > '2015-07-01 09:00:00']
```

```
In [9]: df_complete.shape
```

```
Out[9]: (32038, 9)
```

```
In [10]: # set index
df_complete2 = df_complete.set_index('date_time')
df_complete2.index = pd.DatetimeIndex(df_complete2.index).to_period('H')
```

```
In [11]: # drop variables
target_df = df_complete['traffic_volume']
target_df.head()
```

```
Out[11]: 16166    4273
16167    4469
16168    4625
16169    4462
16170    4996
Name: traffic_volume, dtype: int64
```

## Checking Stationarity with Augmented Dicky-Fuller Test

<https://analyticsindiamag.com/complete-guide-to-dickey-fuller-test-in-time-series-analysis/>

```
In [12]: from statsmodels.tsa.stattools import adfuller
```

```
In [13]: series = target_df.values
```

```
In [14]: # ADF Test
result = adfuller(series, autolag='AIC')
```

```
In [15]: # cite source here: https://analyticsindiamag.com/complete-guide-to-dickey-fuller-test-

print('ADF Statistic: %f' % result[0])

print('p-value: %f' % result[1])

print('Critical Values:')

for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
if result[0] < result[4]["5%"]:
    print ("Reject Ho - Time Series is Stationary")
else:
    print ("Failed to Reject Ho - Time Series is Non-Stationary")
```

```
ADF Statistic: -21.926192
p-value: 0.000000
Critical Values:
    1%: -3.431
    5%: -2.862
   10%: -2.567
Reject Ho - Time Series is Stationary
```


## Determining Parameters (p,d,q)


We know  $d=0$  because the time series is stationary.

```
In [16]: # Finding p (AR Term)
plot_pacf(series)
```

C:\Users\karli\AppData\Roaming\Python\Python39\site-packages\statsmodels\graphics\tsaplo  
ts.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the  
[-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). Y  
ou can use this method now by setting method='ywm'.

```
warnings.warn(
```





## Auto ARIMA

```
In [17]: ! pip install pmdarima --user
```

```
Requirement already satisfied: pmdarima in c:\users\karli\appdata\roaming\python\python39\site-packages (2.0.3)
Requirement already satisfied: Cython!=0.29.18,!0.29.31,>=0.29 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (0.29.24)
Requirement already satisfied: joblib>=0.11 in c:\users\karli\appdata\roaming\python\python39\site-packages (from pmdarima) (1.2.0)
Requirement already satisfied: numpy>=1.21.2 in c:\users\karli\appdata\roaming\python\python39\site-packages (from pmdarima) (1.22.4)
Requirement already satisfied: statsmodels>=0.13.2 in c:\users\karli\appdata\roaming\python\python39\site-packages (from pmdarima) (0.13.5)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (58.0.4)
Requirement already satisfied: pandas>=0.19 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.3.4)
Requirement already satisfied: urllib3 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.26.7)
Requirement already satisfied: scipy>=1.3.2 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.7.1)
Requirement already satisfied: scikit-learn>=0.22 in c:\users\karli\appdata\roaming\python\python39\site-packages (from pmdarima) (1.2.1)
Requirement already satisfied: pytz>=2017.3 in c:\programdata\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2021.3)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\programdata\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2.8.2)
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil>=2.7.3->pandas>=0.19->pmdarima) (1.16.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn>=0.22->pmdarima) (2.2.0)
Requirement already satisfied: packaging>=21.3 in c:\users\karli\appdata\roaming\python\python39\site-packages (from statsmodels>=0.13.2->pmdarima) (23.1)
Requirement already satisfied: patsy>=0.5.2 in c:\programdata\anaconda3\lib\site-packages (from statsmodels>=0.13.2->pmdarima) (0.5.2)
WARNING: Ignoring invalid distribution -quests (c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -umpy (c:\programdata\anaconda3\lib\site-packages)
WARNING: Ignoring invalid distribution -quests (c:\programdata\anaconda3\lib\site-packages)
```

[illegible]


WARNING: Ignoring invalid distribution -equests (c:\programdata\anaconda3\lib\site-packages)

```
In [18]: from pmdarima.arma import auto_arma
```

```
In [19]: # train and test split
train_size = int(len(target_df) * 0.8)
train = target_df[:train_size]
test = target_df[train_size:]
```

```
In [20]: train.plot()
test.plot()
```

Out[20]: <AxesSubplot:>



```
In [21]: arima_model = auto_arma(train, trace=True, information_criterion='bic', max_order = 5)
```

Performing stepwise search to minimize bic

ARIMA(2,0,2)(0,0,0)[0] intercept	: BIC=409019.857, Time=3.69 sec
ARIMA(0,0,0)(0,0,0)[0] intercept	: BIC=461602.504, Time=0.29 sec
ARIMA(1,0,0)(0,0,0)[0] intercept	: BIC=415718.150, Time=0.64 sec
ARIMA(0,0,1)(0,0,0)[0] intercept	: BIC=435734.986, Time=4.36 sec
ARIMA(0,0,0)(0,0,0)[0]	: BIC=495255.962, Time=0.17 sec
ARIMA(1,0,2)(0,0,0)[0] intercept	: BIC=409756.291, Time=2.91 sec
ARIMA(2,0,1)(0,0,0)[0] intercept	: BIC=409026.840, Time=3.11 sec
ARIMA(3,0,2)(0,0,0)[0] intercept	: BIC=408968.620, Time=13.62 sec
ARIMA(3,0,1)(0,0,0)[0] intercept	: BIC=409009.407, Time=6.50 sec
ARIMA(4,0,2)(0,0,0)[0] intercept	: BIC=408904.955, Time=24.20 sec
ARIMA(4,0,1)(0,0,0)[0] intercept	: BIC=408990.616, Time=8.81 sec
ARIMA(5,0,2)(0,0,0)[0] intercept	: BIC=408307.740, Time=39.89 sec
ARIMA(5,0,1)(0,0,0)[0] intercept	: BIC=408794.236, Time=12.31 sec
ARIMA(5,0,3)(0,0,0)[0] intercept	: BIC=408604.551, Time=40.86 sec
ARIMA(4,0,3)(0,0,0)[0] intercept	: BIC=408766.504, Time=30.91 sec
ARIMA(5,0,2)(0,0,0)[0]	: BIC=410777.842, Time=18.99 sec

Best model: ARIMA(5,0,2)(0,0,0)[0] intercept  
Total fit time: 211.328 seconds

```
In [22]: arima_model.summary()
```

Out[22]:

### SARIMAX Results

**Dep. Variable:** y **No. Observations:** 25630  
**Model:** SARIMAX(5, 0, 2) **Log Likelihood** -204108.188  
**Date:** Sun, 07 May 2023 **AIC** 408234.376  
**Time:** 11:52:31 **BIC** 408307.740  
**Sample:** 0 **HQIC** 408258.093  
- 25630  
**Covariance Type:** opg

	coef	std err	z	P> z	[0.025	0.975]
<b>intercept</b>	548.2602	16.349	33.534	0.000	516.216	580.304
<b>ar.L1</b>	1.8325	0.010	187.205	0.000	1.813	1.852
<b>ar.L2</b>	-1.9355	0.019	-101.544	0.000	-1.973	-1.898
<b>ar.L3</b>	1.2753	0.019	67.209	0.000	1.238	1.313
<b>ar.L4</b>	-0.2299	0.012	-19.206	0.000	-0.253	-0.206
<b>ar.L5</b>	-0.1175	0.006	-20.842	0.000	-0.129	-0.106
<b>ma.L1</b>	-0.5065	0.009	-56.699	0.000	-0.524	-0.489
<b>ma.L2</b>	0.8779	0.009	100.832	0.000	0.861	0.895
<b>sigma2</b>	4.883e+05	3010.923	162.182	0.000	4.82e+05	4.94e+05

**Ljung-Box (L1) (Q):** 5.85 **Jarque-Bera (JB):** 39939.19  
**Prob(Q):** 0.02 **Prob(JB):** 0.00  
**Heteroskedasticity (H):** 0.51 **Skew:** 0.73  
**Prob(H) (two-sided):** 0.00 **Kurtosis:** 8.94

Warnings:


[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [28]:

```
test_plot = df_complete[train_size:]
plt.plot(test_plot['date_time'], test)
plt.plot(test_plot['date_time'], arima_model.predict(n_periods=test.shape[0]), color='o')
plt.xticks(rotation=45)
plt.ylabel('Traffic Volume')
plt.show()
```

C:\Users\karli\AppData\Roaming\Python\Python39\site-packages\statsmodels\tsa\base\tsa\_model.py:834: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.  
return get\_prediction\_index(





In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]: